

- `index` — positional label aligned with tool arguments;
- `type` — code or markdown;
- `source` — cell input text;
- `outputs` — optional list of stdout/stderr/display data when scraped;
- `execution_count` — optional prompt number when visible in the [DOM](#).

The host strips or retains execution metadata according to flag `KERNEL_EXECUTION_METADATA_ENABLED`, balancing token budget against diagnostic value for Ask mode error explanation.

3.13.2 Token budgeting

Context packing applies two limits. `MAX_INPUT_TOKENS` caps the full message list; `MAX_NOTEBOOK_CONTEXT_CHARS` caps the notebook slice. When limits are exceeded, older chat history is truncated first; notebook cells are packed from the start of the document unless intent-based packing selects a window around a cited cell index. This policy explains workflow-level Ask behaviour on long notebooks discussed in [Chapter 5](#).

3.14 Security and Trust Model

The trust model assigns capabilities to each component:

- **Extension.** Same-origin [DOM](#) access on Kaggle pages the user has opened; no network egress except native messaging and Kaggle itself.
- **Host.** Holds [API](#) credentials; reads/writes local snapshot files; dispatches tool commands only to the connected extension port.
- **Remote model.** Receives prompts and tool schemas; does not receive browser cookies or Kaggle authentication tokens.

Tool argument validation rejects out-of-range cell indices and malformed JSON before dispatch. Session coercion overwrites model-supplied URLs with the canonical notebook URL from the active turn, reducing the risk of cross-tab misrouting when the model hallucinates a different kernel link.